

Linear Algebra and Applications

Lecture 14-15 Least Squares Method

Zamir Hussain
Department of Mathematics
University of Wah

December 2, 2025

Abstract

These lecture notes introduce the mathematical framework for solving the **least squares** problem, which addresses systems of linear equations ($\mathbf{Ax} = \mathbf{b}$) where an exact solution does not exist. The core objective is defined as finding the vector $\hat{\mathbf{x}}$ that minimizes the norm of the **residual vector** ($\|\mathbf{Ax} - \mathbf{b}\|^2$), thereby finding the closest possible solution. Algebraic and geometric interpretations are provided, along with practical applications in linear regression and data science.

Contents

1	Introduction to Least Squares	3
1.1	The Fundamental Problem	3
1.2	The Residual Concept	3
1.3	Why Squared Norm?	3
2	Geometric Interpretation	3
2.1	Visualizing the Problem	3
2.2	Orthogonality Principle	4
3	Derivation of Normal Equations	4
3.1	From Geometry to Algebra	4
3.2	Step-by-Step Derivation	4
3.3	Properties of Normal Equations	4
4	Existence and Uniqueness of Solutions	5
4.1	Always Existence	5
4.2	Uniqueness Condition	5
4.2.1	Case 1: Unique Solution (Full Column Rank)	5
4.2.2	Case 2: Infinite Solutions (Rank Deficient)	5
4.3	Proof of Optimality (Pythagorean Theorem)	5

5	Application: Linear Regression	6
5.1	Problem Setup	6
5.2	Matrix Formulation	6
5.3	Step-by-Step Example	6
6	Practical Applications in Data Science & ML	7
6.1	Linear Regression (Foundation of ML)	7
6.2	Recommender Systems	7
6.3	Computer Vision	7
6.4	Natural Language Processing	8
6.5	Signal Processing	8
7	Numerical Considerations	8
7.1	Conditioning Issues	8
7.2	Better Numerical Methods	8
8	Practice Problems	8
8.1	Problem 1: Basic Least Squares	8
8.2	Problem 2: Polynomial Fitting	9
8.3	Problem 3: Rank Deficiency	10
9	Summary & Key Takeaways	11
9.1	Key Formulas	11
9.2	Geometric Insight	11
9.3	Practical Advice	11
9.4	Applications Checklist	12
10	Additional Resources	12
10.1	Textbooks:	12
10.2	Online Courses:	12
10.3	Software Libraries:	12

1 Introduction to Least Squares

1.1 The Fundamental Problem

In many practical situations, we encounter systems of linear equations where the number of equations exceeds the number of unknowns:

$$\mathbf{Ax} = \mathbf{b}$$

where:

- $\mathbf{A} \in \mathbb{R}^{m \times n}$ is an $m \times n$ matrix
- $\mathbf{x} \in \mathbb{R}^n$ is the unknown vector
- $\mathbf{b} \in \mathbb{R}^m$ is the observation vector
- Typically $m > n$ (more equations than unknowns)

Definition

The Challenge: When $\mathbf{b} \notin R(\mathbf{A})$ (i.e., $\mathbf{b}$ is not in the column space of $\mathbf{A}$), there is **no exact solution**.

1.2 The Residual Concept

We define the **residual vector** as:

$$\mathbf{r}(\mathbf{x}) = \mathbf{Ax} - \mathbf{b}$$

The magnitude of this residual measures how "good" our solution is. When no exact solution exists, we seek the "best possible" solution.

Definition

Least Squares Objective: Find $\hat{\mathbf{x}}$ that minimizes the squared Euclidean norm of the residual:

$$\min_{\mathbf{x} \in \mathbb{R}^n} \|\mathbf{Ax} - \mathbf{b}\|^2$$

1.3 Why Squared Norm?

- **Differentiable everywhere** (unlike the absolute value)
- **Emphasizes larger errors** more heavily
- **Leads to nice analytical solutions**
- Equivalent to minimizing $(\mathbf{Ax} - \mathbf{b})^T(\mathbf{Ax} - \mathbf{b})$

2 Geometric Interpretation

2.1 Visualizing the Problem

Consider a simple case with $m = 3$ and $n = 2$:

- The columns of $\mathbf{A}$ span a plane in $\mathbb{R}^3$ (the column space $R(\mathbf{A})$)

- Vector $\mathbf{b}$ is typically not in this plane
- We want to find the point in $R(\mathbf{A})$ closest to $\mathbf{b}$

Geometric interpretation:
 $R(\mathbf{A})$ is a subspace (plane through origin)
 $\mathbf{b}$ lies outside this plane
 Find $\mathbf{A}\hat{\mathbf{x}} \in R(\mathbf{A})$ closest to $\mathbf{b}$

2.2 Orthogonality Principle

Theorem

Theorem (Orthogonality Principle): The optimal residual $\hat{\mathbf{r}} = \mathbf{A}\hat{\mathbf{x}} - \mathbf{b}$ must be **orthogonal** to every vector in $R(\mathbf{A})$.
 Mathematically: $\hat{\mathbf{r}} \perp R(\mathbf{A})$ or $\hat{\mathbf{r}} \in R(\mathbf{A})^\perp$

Why? If the residual weren't orthogonal, we could find a point in $R(\mathbf{A})$ even closer to $\mathbf{b}$.

3 Derivation of Normal Equations

3.1 From Geometry to Algebra

The orthogonality condition can be expressed mathematically:

$$\mathbf{w}^T(\mathbf{A}\hat{\mathbf{x}} - \mathbf{b}) = 0 \quad \text{for all } \mathbf{w} \in R(\mathbf{A})$$

Since any $\mathbf{w} \in R(\mathbf{A})$ can be written as $\mathbf{A}\mathbf{v}$ for some $\mathbf{v}$, we have:

$$(\mathbf{A}\mathbf{v})^T(\mathbf{A}\hat{\mathbf{x}} - \mathbf{b}) = 0 \quad \text{for all } \mathbf{v}$$

3.2 Step-by-Step Derivation

1. **Start with orthogonality condition:**

$$(\mathbf{A}\mathbf{v})^T(\mathbf{A}\hat{\mathbf{x}} - \mathbf{b}) = 0 \quad \forall \mathbf{v}$$

2. **Apply transpose property:**

$$\mathbf{v}^T \mathbf{A}^T (\mathbf{A}\hat{\mathbf{x}} - \mathbf{b}) = 0 \quad \forall \mathbf{v}$$

3. **For this to hold for all $\mathbf{v}$:**

$$\mathbf{A}^T (\mathbf{A}\hat{\mathbf{x}} - \mathbf{b}) = \mathbf{0}$$

4. **Rearrange to get Normal Equations:**

$$\mathbf{A}^T \mathbf{A} \hat{\mathbf{x}} = \mathbf{A}^T \mathbf{b}$$

3.3 Properties of Normal Equations

- **Size:** The system $\mathbf{A}^T \mathbf{A} \hat{\mathbf{x}} = \mathbf{A}^T \mathbf{b}$ is $n \times n$ (smaller than original $m \times n$)
- **Symmetry:** $\mathbf{A}^T \mathbf{A}$ is always symmetric
- **Positive Semi-definite:** $\mathbf{A}^T \mathbf{A}$ is always positive semi-definite
- **Always solvable:** $\mathbf{A}^T \mathbf{b}$ is always in the column space of $\mathbf{A}^T \mathbf{A}$

4 Existence and Uniqueness of Solutions

4.1 Always Existence

Key Concept

Key Fact: The normal equations $A^T A \hat{x} = A^T b$ **always have at least one solution**, regardless of A .

Why? Because $A^T b$ is always in the column space of $A^T A$.

4.2 Uniqueness Condition

The uniqueness of the solution depends on the **rank** of matrix A :

4.2.1 Case 1: Unique Solution (Full Column Rank)

If A has **linearly independent columns** ($\text{rank} = n$), then:

- $A^T A$ is **invertible**
- The unique solution is:

$$\hat{x} = (A^T A)^{-1} A^T b$$

- The matrix $(A^T A)^{-1} A^T$ is called the **Moore-Penrose pseudo-inverse** of A

4.2.2 Case 2: Infinite Solutions (Rank Deficient)

If A has **linearly dependent columns** ($\text{rank} < n$), then:

- $A^T A$ is **singular** (not invertible)
- There are infinitely many solutions
- All solutions give the **same minimum residual**
- The general solution is:

$$\hat{x} = \hat{x}_p + w, \quad w \in N(A)$$

where $\hat{x}_p$ is any particular solution and $N(A)$ is the null space of A

4.3 Proof of Optimality (Pythagorean Theorem)

Theorem

Theorem (Optimality): $\hat{x}$ minimizes $\|Ax - b\|^2$ **if and only if** it satisfies $A^T A \hat{x} = A^T b$.

Proof: Let $\hat{x}$ satisfy the normal equations, and let x be any other vector.

1. Compute the residual for x :

$$r = Ax - b$$

2. Write r in terms of $\hat{r}$:

$$r = (A\hat{x} - b) + A(x - \hat{x}) = \hat{r} + A(x - \hat{x})$$

3. Compute the squared norm:

$$\|\mathbf{r}\|^2 = \|\hat{\mathbf{r}} + \mathbf{A}(\mathbf{x} - \hat{\mathbf{x}})\|^2$$

4. Expand using dot product properties:

$$\|\mathbf{r}\|^2 = \|\hat{\mathbf{r}}\|^2 + 2\hat{\mathbf{r}}^T \mathbf{A}(\mathbf{x} - \hat{\mathbf{x}}) + \|\mathbf{A}(\mathbf{x} - \hat{\mathbf{x}})\|^2$$

5. The cross term vanishes because $\hat{\mathbf{r}}$ is orthogonal to $R(\mathbf{A})$:

$$\hat{\mathbf{r}}^T \mathbf{A}(\mathbf{x} - \hat{\mathbf{x}}) = (\mathbf{A}^T \hat{\mathbf{r}})^T (\mathbf{x} - \hat{\mathbf{x}}) = \mathbf{0}^T (\mathbf{x} - \hat{\mathbf{x}}) = 0$$

6. Therefore:

$$\|\mathbf{r}\|^2 = \|\hat{\mathbf{r}}\|^2 + \|\mathbf{A}(\mathbf{x} - \hat{\mathbf{x}})\|^2 \geq \|\hat{\mathbf{r}}\|^2$$

This proves $\hat{\mathbf{x}}$ minimizes the residual!

5 Application: Linear Regression

5.1 Problem Setup

Scenario: We want to fit a straight line to data points (x_i, y_i) , $i = 1, \dots, m$.

Model: $y = \beta_0 + \beta_1 x$ (intercept β_0 , slope β_1)

For each data point: $y_i \approx \beta_0 + \beta_1 x_i$

5.2 Matrix Formulation

We can write this as $\mathbf{A}\boldsymbol{\beta} \approx \mathbf{b}$:

$$\underbrace{\begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_m \end{bmatrix}}_{\mathbf{A}} \underbrace{\begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}}_{\boldsymbol{\beta}} \approx \underbrace{\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}}_{\mathbf{b}}$$

5.3 Step-by-Step Example

Data: Points (1, 1), (2, 3), (3, 3), (4, 5)

1. **Construct matrices:**

$$\mathbf{A} = \begin{bmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \\ 1 & 4 \end{bmatrix}, \quad \mathbf{b} = \begin{bmatrix} 1 \\ 3 \\ 3 \\ 5 \end{bmatrix}$$

2. **Compute $\mathbf{A}^T \mathbf{A}$:**

$$\mathbf{A}^T \mathbf{A} = \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \end{bmatrix} \begin{bmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \\ 1 & 4 \end{bmatrix} = \begin{bmatrix} 4 & 10 \\ 10 & 30 \end{bmatrix}$$

3. Compute $\mathbf{A}^T \mathbf{b}$:

$$\mathbf{A}^T \mathbf{b} = \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \end{bmatrix} \begin{bmatrix} 1 \\ 3 \\ 3 \\ 5 \end{bmatrix} = \begin{bmatrix} 12 \\ 36 \end{bmatrix}$$

4. Solve normal equations:

$$\begin{bmatrix} 4 & 10 \\ 10 & 30 \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix} = \begin{bmatrix} 12 \\ 36 \end{bmatrix}$$

5. Solve the system:

$$\begin{cases} 4\beta_0 + 10\beta_1 = 12 \\ 10\beta_0 + 30\beta_1 = 36 \end{cases}$$

Multiply first equation by 2.5: $10\beta_0 + 25\beta_1 = 30$

Subtract from second equation: $5\beta_1 = 6 \Rightarrow \beta_1 = 1.2$

Substitute: $4\beta_0 + 10(1.2) = 12 \Rightarrow 4\beta_0 = 0 \Rightarrow \beta_0 = 0$

6. **Solution:** $\beta_0 = 0, \beta_1 = 1.2$

Fitted line: $y = 1.2x$

7. Verify:

$$\mathbf{A}\hat{\boldsymbol{\beta}} = \begin{bmatrix} 1.2 \\ 2.4 \\ 3.6 \\ 4.8 \end{bmatrix}, \quad \hat{\mathbf{r}} = \mathbf{A}\hat{\boldsymbol{\beta}} - \mathbf{b} = \begin{bmatrix} 0.2 \\ -0.6 \\ 0.6 \\ -0.2 \end{bmatrix}$$

$$\|\hat{\mathbf{r}}\|^2 = 0.04 + 0.36 + 0.36 + 0.04 = 0.8$$

6 Practical Applications in Data Science & ML

6.1 Linear Regression (Foundation of ML)

- **Ordinary Least Squares (OLS):** Basic linear regression
- **Ridge Regression:** $\min \|\mathbf{Ax} - \mathbf{b}\|^2 + \lambda \|\mathbf{x}\|^2$ (adds L2 regularization)
- **Lasso Regression:** $\min \|\mathbf{Ax} - \mathbf{b}\|^2 + \lambda \|\mathbf{x}\|_1$ (adds L1 regularization, promotes sparsity)

6.2 Recommender Systems

- **Collaborative Filtering:** Predict user ratings for items
- **Matrix Factorization:** $\min \|\mathbf{R} - \mathbf{UV}^T\|^2$ where $\mathbf{R}$ is rating matrix
- **Netflix Prize:** Famous competition won using matrix factorization techniques

6.3 Computer Vision

- **Image Reconstruction:** Reconstruct images from partial data
- **Structure from Motion:** Reconstruct 3D structure from 2D images
- **Camera Calibration:** Estimate camera parameters from known points

6.4 Natural Language Processing

- **Latent Semantic Analysis (LSA):** Reduce dimensionality of term-document matrices
- **Word Embeddings:** Methods like GloVe use matrix factorization

6.5 Signal Processing

- **Compressive Sensing:** Recover signals from fewer measurements
- **System Identification:** Estimate parameters of dynamical systems
- **Filter Design:** Design digital filters with desired frequency response

7 Numerical Considerations

7.1 Conditioning Issues

The matrix $\mathbf{A}^T \mathbf{A}$ can be **ill-conditioned** even when $\mathbf{A}$ is reasonably conditioned.

Condition number: $\kappa(\mathbf{A}^T \mathbf{A}) = [\kappa(\mathbf{A})]^2$

Example: If $\kappa(\mathbf{A}) = 10^3$, then $\kappa(\mathbf{A}^T \mathbf{A}) = 10^6$!

7.2 Better Numerical Methods

Instead of solving normal equations directly, use:

1. **QR Factorization:** $\mathbf{A} = \mathbf{Q}\mathbf{R}$, then solve $\mathbf{R}\mathbf{x} = \mathbf{Q}^T \mathbf{b}$
2. **Singular Value Decomposition (SVD):** $\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$, then $\hat{\mathbf{x}} = \mathbf{V}\mathbf{\Sigma}^{-1}\mathbf{U}^T \mathbf{b}$
3. **Iterative Methods:** For very large sparse systems (conjugate gradient)

8 Practice Problems

8.1 Problem 1: Basic Least Squares

Given the system:

$$\begin{cases} x_1 + x_2 = 2 \\ 2x_1 + x_2 = 3 \\ 3x_1 + x_2 = 5 \end{cases}$$

- (a) Write in matrix form $\mathbf{A}\mathbf{x} = \mathbf{b}$
- (b) Find the least squares solution
- (c) Compute the residual norm

Solution 1

(a) Matrix form:

$$\mathbf{A} = \begin{bmatrix} 1 & 1 \\ 2 & 1 \\ 3 & 1 \end{bmatrix}, \quad \mathbf{b} = \begin{bmatrix} 2 \\ 3 \\ 5 \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

(b) Compute normal equations:

$$\mathbf{A}^T \mathbf{A} = \begin{bmatrix} 1 & 2 & 3 \\ 1 & 1 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 \\ 2 & 1 \\ 3 & 1 \end{bmatrix} = \begin{bmatrix} 14 & 6 \\ 6 & 3 \end{bmatrix}$$

$$\mathbf{A}^T \mathbf{b} = \begin{bmatrix} 1 & 2 & 3 \\ 1 & 1 & 1 \end{bmatrix} \begin{bmatrix} 2 \\ 3 \\ 5 \end{bmatrix} = \begin{bmatrix} 23 \\ 10 \end{bmatrix}$$

Solve:

$$\begin{cases} 14x_1 + 6x_2 = 23 \\ 6x_1 + 3x_2 = 10 \end{cases}$$

Multiply second equation by 2: $12x_1 + 6x_2 = 20$

Subtract from first: $2x_1 = 3 \Rightarrow x_1 = 1.5$

Substitute: $6(1.5) + 3x_2 = 10 \Rightarrow 9 + 3x_2 = 10 \Rightarrow x_2 = \frac{1}{3}$

Solution: $x_1 = 1.5, x_2 = \frac{1}{3}$

(c) Residual:

$$\mathbf{r} = \mathbf{Ax} - \mathbf{b} = \begin{bmatrix} 1.5 + \frac{1}{3} - 2 \\ 3 + \frac{1}{3} - 3 \\ 4.5 + \frac{1}{3} - 5 \end{bmatrix} = \begin{bmatrix} -\frac{1}{6} \\ \frac{1}{3} \\ -\frac{1}{6} \end{bmatrix}$$

$$\|\mathbf{r}\|^2 = \frac{1}{36} + \frac{1}{9} + \frac{1}{36} = \frac{1}{36} + \frac{4}{36} + \frac{1}{36} = \frac{6}{36} = \frac{1}{6}$$

8.2 Problem 2: Polynomial Fitting

Fit a quadratic polynomial $y = a + bx + cx^2$ to the points: $(-1, 2), (0, 1), (1, 3), (2, 5)$

Solution 2

Model: $y = a + bx + cx^2$

For each point:

- $(-1, 2)$: $a - b + c = 2$
- $(0, 1)$: $a = 1$
- $(1, 3)$: $a + b + c = 3$
- $(2, 5)$: $a + 2b + 4c = 5$

Matrix form:

$$\mathbf{A} = \begin{bmatrix} 1 & -1 & 1 \\ 1 & 0 & 0 \\ 1 & 1 & 1 \\ 1 & 2 & 4 \end{bmatrix}, \quad \mathbf{b} = \begin{bmatrix} 2 \\ 1 \\ 3 \\ 5 \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} a \\ b \\ c \end{bmatrix}$$

Normal equations:

$$\mathbf{A}^T \mathbf{A} = \begin{bmatrix} 1 & 1 & 1 & 1 \\ -1 & 0 & 1 & 2 \\ 1 & 0 & 1 & 4 \end{bmatrix} \begin{bmatrix} 1 & -1 & 1 \\ 1 & 0 & 0 \\ 1 & 1 & 1 \\ 1 & 2 & 4 \end{bmatrix} = \begin{bmatrix} 4 & 2 & 6 \\ 2 & 6 & 8 \\ 6 & 8 & 18 \end{bmatrix}$$

$$\mathbf{A}^T \mathbf{b} = \begin{bmatrix} 1 & 1 & 1 & 1 \\ -1 & 0 & 1 & 2 \\ 1 & 0 & 1 & 4 \end{bmatrix} \begin{bmatrix} 2 \\ 1 \\ 3 \\ 5 \end{bmatrix} = \begin{bmatrix} 11 \\ 11 \\ 31 \end{bmatrix}$$

Solve:

$$\begin{cases} 4a + 2b + 6c = 11 \\ 2a + 6b + 8c = 11 \\ 6a + 8b + 18c = 31 \end{cases}$$

Solving systematically gives: $a = -0.05$, $b = -1.15$, $c = 2.25$

Polynomial: $y = -0.05 - 1.15x + 2.25x^2$

8.3 Problem 3: Rank Deficiency

Consider:

$$\mathbf{A} = \begin{bmatrix} 1 & 2 & 3 \\ 1 & 2 & 3 \\ 1 & 2 & 3 \end{bmatrix}, \quad \mathbf{b} = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$$

- What is the rank of $\mathbf{A}$?
- Show that the normal equations have infinitely many solutions
- Find the minimum residual norm

Solution 3

(a) Rank: All rows are identical, so $\text{rank}(\mathbf{A}) = 1$

(b) Normal equations:

$$\mathbf{A}^T \mathbf{A} = \begin{bmatrix} 1 & 1 & 1 \\ 2 & 2 & 2 \\ 3 & 3 & 3 \end{bmatrix} \begin{bmatrix} 1 & 2 & 3 \\ 1 & 2 & 3 \\ 1 & 2 & 3 \end{bmatrix} = \begin{bmatrix} 3 & 6 & 9 \\ 6 & 12 & 18 \\ 9 & 18 & 27 \end{bmatrix}$$

$$\mathbf{A}^T \mathbf{b} = \begin{bmatrix} 1 & 1 & 1 \\ 2 & 2 & 2 \\ 3 & 3 & 3 \end{bmatrix} \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix} = \begin{bmatrix} 6 \\ 12 \\ 18 \end{bmatrix}$$

All equations are multiples of: $x_1 + 2x_2 + 3x_3 = 2$

This is a plane in $\mathbb{R}^3$ - infinitely many solutions!

(c) Minimum residual: Any solution satisfying $x_1 + 2x_2 + 3x_3 = 2$ gives:

$$\mathbf{Ax} = \begin{bmatrix} 2 \\ 2 \\ 2 \end{bmatrix}, \quad \mathbf{r} = \mathbf{Ax} - \mathbf{b} = \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix}$$

$$\|\mathbf{r}\|^2 = 1^2 + 0^2 + (-1)^2 = 2$$

All solutions give the same residual norm of 2.

9 Summary & Key Takeaways

9.1 Key Formulas

1. **Least Squares Problem:** $\min_{\mathbf{x}} \|\mathbf{Ax} - \mathbf{b}\|^2$
2. **Normal Equations:** $\mathbf{A}^T \mathbf{A} \hat{\mathbf{x}} = \mathbf{A}^T \mathbf{b}$
3. **Unique Solution:** $\hat{\mathbf{x}} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{b}$ (if $\mathbf{A}$ has full column rank)
4. **Residual:** $\hat{\mathbf{r}} = \mathbf{A} \hat{\mathbf{x}} - \mathbf{b}$

9.2 Geometric Insight

- Least squares finds the **orthogonal projection** of $\mathbf{b}$ onto $R(\mathbf{A})$
- The optimal residual is orthogonal to the column space
- The solution minimizes the Euclidean distance from $\mathbf{b}$ to $R(\mathbf{A})$

9.3 Practical Advice

1. **Check rank** of $\mathbf{A}$ before solving
2. **Avoid directly computing** $(\mathbf{A}^T \mathbf{A})^{-1}$ for numerical stability
3. Use **QR factorization** or **SVD** for better numerical solutions
4. Consider **regularization** (Ridge/Lasso) if $\mathbf{A}^T \mathbf{A}$ is ill-conditioned

9.4 Applications Checklist

- ✓ Linear regression and curve fitting
- ✓ System identification and parameter estimation
- ✓ Signal processing and filtering
- ✓ Machine learning (basis of many algorithms)
- ✓ Computer vision and image processing
- ✓ Recommender systems and collaborative filtering

10 Additional Resources

10.1 Textbooks:

- Strang, G. *Linear Algebra and Its Applications*
- Golub, G. H., & Van Loan, C. F. *Matrix Computations*
- Hastie, T., Tibshirani, R., & Friedman, J. *The Elements of Statistical Learning*

10.2 Online Courses:

- MIT OpenCourseWare: Linear Algebra (18.06)
- Coursera: Machine Learning by Andrew Ng
- edX: Linear Algebra - Foundations to Frontiers

10.3 Software Libraries:

- Python: `numpy.linalg.lstsq`, `scipy.linalg.lstsq`
- MATLAB: `\` operator (backslash) or `mldivide`
- R: `lm()` function for linear models

Remember: Least squares is more than just a mathematical trick—it's a fundamental principle for extracting meaningful information from noisy data, with applications spanning science, engineering, and data analysis.